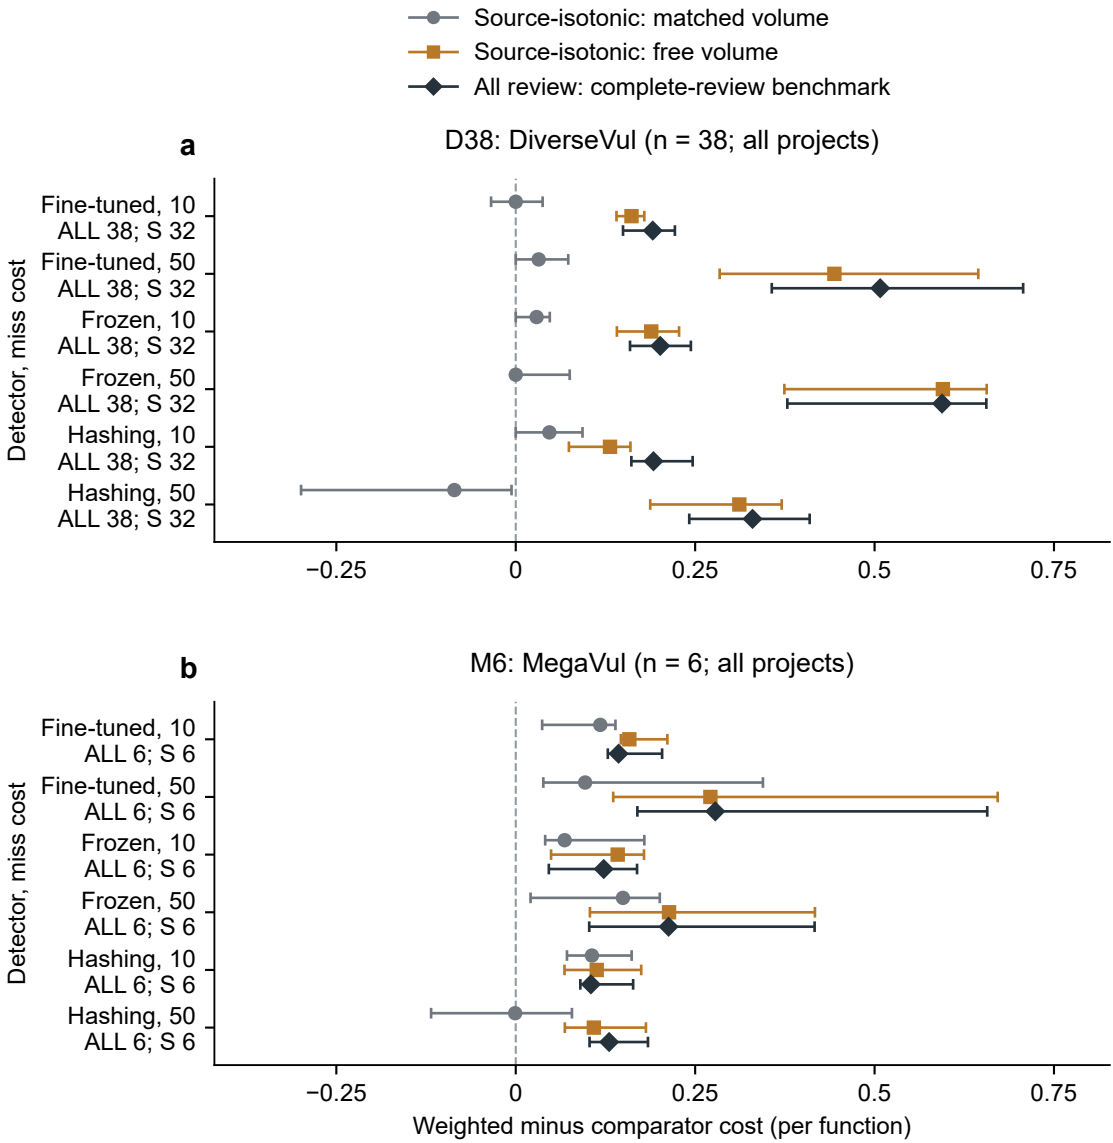

Median paired cost difference; 95% project bootstrap CI (10,000 draws).
ALL = all projects; S = every actual domain fit available; no supported-only filtering.
Miss costs 10/50; false-alert cost 1; review cost 0.1; correct review assumed.
Negative differences favor weighted. Domain-failure review fallback is retained.